



# 703308 VO High-Performance Computing MPI Groups, Communicators and One-Sided Communication

Philipp Gschwandtner

# Overview

---

- ▶ **communicators and groups**
  - ▶ more ways of limiting and controlling collective communication
- ▶ **one-sided communication**
  - ▶ decouples data access and synchronization
- ▶ **error handling**

# Motivation




---

- ▶ Real-world applications are rarely a single component
  - ▶ often MPMD
  - ▶ usually combination of libraries (e.g. molecular dynamics and quantum mechanics)
- ▶ Adds several new complexities compared to single-component software
  - ▶ collective communication via `MPI_COMM_WORLD`?
  - ▶ how to identify sub-programs and communicate between and within them?

## Motivation cont'd


---

- ▶ unstructured grid with cell and face element types
  - ▶ assume both require global communication (e.g. reduction) among their types
  
- ▶ a communicator per element type would be useful
  - ▶ but how?



## Motivation cont'd

- ▶ What about NUMA?
  - ▶ virtual topologies do not reflect e.g. shared memory address spaces
- ▶ consider shared memory node with 4 CPUs and many cores per CPU
- ▶ local collective communication among cores of a node is cheap
  - ▶ how to limit?
  - ▶ construct a communicator per node?
  - ▶ do this in hardware- and compiler-independent way?



# Communicators and groups

---

- ▶ Communicators and groups hold sets of ranks
  - ▶ directly used for e.g. collective communication
  - ▶ also required for identifying single ranks
  - ▶ remember MPI basics lecture: everything in MPI is relative to a communicator or group
- ▶ Why not stick to **MPI\_COMM\_WORLD**?
  - ▶ isolate application sub-programs
    - ▶ individual processing steps running in parallel (task parallelism, MPMD)
    - ▶ domain decomposition (SPMD, c.f. slicing Cartesian topologies)
    - ▶ libraries (portability)
  - ▶ usability
    - ▶ add user-defined attributes such as topologies
  - ▶ performance
    - ▶ re-numbering of ranks (virtual topology vs. hardware topology)

# Communicators and groups cont'd

---

## ▶ MPI\_Group

- ▶ holds ordered set of ranks
- ▶ ordering is given by mapping process identifier (e.g. PID) to rank number
- ▶ construction of and operations on groups are always **local** operations

## ▶ MPI\_Comm

- ▶ holds an MPI\_Group
  - ▶ transitively holds ordered set of ranks
- ▶ can hold attributes (e.g. topology)
- ▶ constructed from groups
- ▶ construction of communicators are **non-local** operations (collectives)

# Operations on MPI\_Group

---

## ▶ Constructors

- ▶ `MPI_Comm_group(...)`
- ▶ `MPI_Group_union(...)`
- ▶ `MPI_Group_intersection(...)`
- ▶ `MPI_Group_difference(...)`
- ▶ `MPI_Group_incl(...)`
- ▶ `MPI_Group_excl(...)`
- ▶ `MPI_Group_range_incl(...)`
- ▶ `MPI_Group_range_excl(...)`

## ▶ Accessors

- ▶ `MPI_Group_size(...)`
- ▶ `MPI_Group_rank(...)`
- ▶ `MPI_Group_compare(...)`
  - ▶ result is `MPI_IDENT`, `MPI_SIMILAR` or `MPI_UNEQUAL`

## ▶ Destructor

- ▶ `MPI_Group_free(...)`

# Operations on MPI\_Comm

---

- ▶ Constructors

- ▶ `MPI_Comm_dup(...)`
- ▶ `MPI_Comm_create(...)`
- ▶ `MPI_Comm_split(...)`
- ▶ Convenience constructors such as  
`MPI_Cart_sub(...)`

- ▶ Accessors


- ▶ `MPI_Comm_size(...)`
- ▶ `MPI_Comm_rank(...)`
- ▶ `MPI_Comm_compare(...)`
  - ▶ result is `MPI_IDENT`, `MPI_SIMILAR`,  
`MPI_CONGRUENT` or `MPI_UNEQUAL`

- ▶ Destructor

- ▶ `MPI_Comm_free(...)`

# Group and communicator workflow

- ▶ start with MPI\_COMM\_WORLD
- ▶ construct group(s) of rank subsets and modify as required
  - ▶ MPI\_Group\_union(),  
MPI\_Group\_range\_incl(), ...
- ▶ create new communicator from group and use for communication




## Splitting communicators

---

- ▶ `int MPI_Comm_split(MPI_Comm comm, int color, int key, MPI_Comm* newcomm)`
  - ▶ `comm`: current communicator
  - ▶ `color`: control of subset assignment (same color: same new communicator)
  - ▶ `key`: control of rank assignment (0: sorted as in `comm`; otherwise according to ascending key values)
  - ▶ `newcomm`: new communicator
- ▶ `MPI_Comm_split_type(...)`
  - ▶ allows to split dependent on hardware properties

# MPI\_Comm\_split example

```
MPI_Comm newComm;  
MPI_Comm_rank(MPI_COMM_WORLD, &myRank);  
int myColor = myRank / 4;  
MPI_Comm_split(MPI_COMM_WORLD, myColor,  
    myRank, &newComm);  
MPI_Comm_rank(newComm, &newRank);
```



# Solutions to motivation examples


---

```
MPI_Comm newComm;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
int color =
    (elementType == TYPE_FACES);
MPI_Comm_split(MPI_COMM_WORLD,
    color, rank, &newComm);
```

```
MPI_Comm newComm;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_split_type(MPI_COMM_WORLD,
    MPI_COMM_TYPE_SHARED, rank,
    MPI_INFO_NULL, &newComm);
// also: OMPI_COMM_TYPE_CORE,
// OMPI_COMM_TYPE_L1CACHE,
// OMPI_COMM_TYPE_L2CACHE,
// OMPI_COMM_TYPE_L3CACHE, ...
```

# Intra- and inter-communicators

- ▶ intra-communicator
  - ▶ collection of ranks that can send messages to each other via point-to-point and collectives
  - ▶ e.g. MPI\_COMM\_WORLD
- ▶ inter-communicator
  - ▶ collection of ranks from disjoint intra-communicators
  - ▶ allows sending messages between communicators





## One-sided Communication



# Motivation

---

- ▶ **message-passing paradigm**
  - ▶ fits distributed memory systems well
  - ▶ data transfers among distinct address spaces require network communication
  - ▶ requires explicit communication
  - ▶ downside: little control over message & synchronization aggregation
- ▶ **shared memory paradigm**
  - ▶ no message passing required
  - ▶ data transfer aggregation possible – write multiple bytes, elements, ... in one go
  - ▶ much more convenient from a user and performance perspective
    - ▶ does not necessarily require receiving side to participate
  - ▶ also, messages are needless overhead on shared-memory systems
    - ▶ send/recv function call, management, etc., just for a memcpy in the same address space?


## MPI's solution: one-sided communication

---

- ▶ classic point-to-point (“*two-side*”) communication implies synchronization
  - ▶ at every data transfer action
  - ▶ incurs a lot of overhead in the presence of many messages
- ▶ one-sided communication decouples data movement and synchronization
  - ▶ ranks expose a “*window*” of rank-local memory
  - ▶ can be accessed by other ranks using remote memory access (RMA)
  - ▶ data accesses do not necessarily require action on the rank exposing memory
  - ▶ both read and write are possible
    - ▶ ranks no longer identify as “*sender*” and “*receiver*” but as “*origin*” and “*target*” instead


# One-sided communication workflow

- ▶ allocate buffer and window
  - ▶ can ask MPI to allocate fast memory
- ▶ open window (“*start epoch*”)
  - ▶ synchronization point
  - ▶ allows data access by remote ranks
- ▶ close window (“*end epoch*”)
  - ▶ synchronization point
  - ▶ commits data accesses
- ▶ deallocate window and buffer




## MPI one-sided communication cont'd

### ▶ two-sided



### ▶ one-sided




# Means of synchronization

---

- ▶ active target synchronization
  - ▶ target participates in synchronization
  - ▶ similar to message-passing paradigm
  - ▶ Uses either `MPI_Win_fence()` or “post-start-complete-wait”
  - ▶ often preferred for bulk-synchronous parallel programs
    - ▶ all ranks execute computation and communication steps more or less in sync
    - ▶ e.g. structured grid with ghost cell exchange
- ▶ passive target synchronization
  - ▶ target does not synchronize
  - ▶ similar to shared memory paradigm
  - ▶ uses `MPI_Win_lock()` and `MPI_Win_unlock()`
  - ▶ often preferred for dynamic, independent access patterns
    - ▶ e.g. irregular codes


# Active target synchronization: fence

- ▶ collective synchronization
  - ▶ origin/target not specified
- ▶ all control the epoch
  - ▶ starts/ends all epochs on all participating ranks
- ▶ fence enforces synchronization




# Active target synchronization: post/start/complete/wait

- ▶ selective synchronization
  - ▶ origin and target specify a group they communicate with
- ▶ both control their epochs
  - ▶ origin: start/complete
  - ▶ target: post/wait
- ▶ synchronization calls may block to enforce ordering



# Passive target synchronization: lock/unlock

- ▶ target neither involved in data transfer, nor in synchronization
- ▶ origin has full control over epoch
- ▶ resembles shared memory programming models (e.g. Pthreads, std::mutex, ...)
  - ▶ but not the same
  - ▶ no critical section!




# Implications of one-sided communication

---

- ▶ several benefits
  - ▶ allows dynamic access patterns (e.g. when target rank does not know number and ranks of origins; or have two ranks communicate data through a third rank who does not participate)
  - ▶ reduce synchronization overhead for multiple data transfers
  - ▶ reduce management overhead on receiver side (e.g. tag matching)
  - ▶ performance gain
  - ▶ reduce coding effort on receiver side
- ▶ drawbacks
  - ▶ no send/receive matching
  - ▶ operations are not explicitly visible on the receiver side
  - ▶ user responsible for correct order of reads/writes (race conditions)
  - ▶ only non-blocking communication

# Performance comparison

- ▶ LCC2, openmpi/3.1.1  
2 ranks, one per node
- ▶ rank 0 sends  $10^8$  int to rank 1
  - ▶ once using  $10^8$  plain send/recv calls
  - ▶ once using  $10^8$  MPI\_Put() calls with MPI\_Fence() synchronization
- ▶ execution time reduced by 2.26x
  - ▶ only between two ranks
  - ▶ but an edge case stress test



## Optional window fence assertions

---

- ▶ **MPI\_MODE\_NOSTORE**
  - ▶ local window was not updated by local store, local get or receive calls since last fence
- ▶ **MPI\_MODE\_NOPUT**
  - ▶ local window will not be updated by put or accumulate until next fence
- ▶ **MPI\_MODE\_NOPRECEDE**
  - ▶ fence does not complete any sequence of locally issued RMA calls
- ▶ **MPI\_MODE\_NOSUCCEED**
  - ▶ fence does not start any sequence of locally issued RMA calls
- ▶ none of these are required, but they can improve performance

## Four window models

---

- ▶ **MPI\_Win\_create(...)**
  - ▶ private memory buffer already allocated, use as window
- ▶ **MPI\_Win\_allocate(...)**
  - ▶ allocate buffer and use as window
- ▶ **MPI\_Win\_create\_dynamic(...)**
  - ▶ expose a buffer which is not available yet
  - ▶ use later with `MPI_Win_attach()`/`MPI_Win_detach()`
- ▶ **MPI\_Win\_allocate\_shared(...)**
  - ▶ allocate and use buffer in shared memory segment of the OS
  - ▶ only works for `MPI_COMM_TYPE_SHARED`

# Transferring data: put & get

---

- ▶ `int MPI_Put(const void* origin_addr, int origin_count, MPI_Datatype origin_datatype, int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype, MPI_Win win)`
  - ▶ `origin_addr`: local address of data to put (aka “send buffer”)
  - ▶ `origin_count`: number of elements on origin side
  - ▶ `origin_datatype`: type of elements on origin side
  - ▶ `target_rank`: target rank
  - ▶ `target_disp`: offset of target address to base address of target window
  - ▶ `target_count`: number of elements on target side
  - ▶ `target_datatype`: type of elements on target side
  - ▶ `win`: window handle
- ▶ **`MPI_Get(...)`**
  - ▶ transfer data from target to origin

## Transferring data: accumulate

---

- ▶ **MPI\_Accumulate(...)**
  - ▶ transfer data from origin and accumulate atomically at target
  - ▶ can only use predefined operations of **MPI\_Reduce(...)** (e.g. **MPI\_SUM**)
  - ▶ use **MPI\_REPLACE** to get atomic put
- ▶ **MPI\_Get\_accumulate(...)**
  - ▶ same as **MPI\_Accumulate(...)** but store target buffer data in result buffer before accumulating
  - ▶ use with **MPI\_NO\_OP** to implement atomic get or **MPI\_REPLACE** for atomic swap

# Transferring data: single-element atomics

---

- ▶ **`MPI_Compare_and_swap(...)`**
  - ▶ atomic swap if data at target buffer matches comparison value
  - ▶ must be single element
  - ▶ must be predefined integer, logical or byte type
  
- ▶ **`MPI_Fetch_and_op(...)`**
  - ▶ variant of `MPI_Get_accumulate(...)`, available for hardware optimization
  - ▶ implements a subset of `MPI_Get_accumulate(...)`'s generic functionality
    - ▶ must be single-element
    - ▶ must be predefined data type

# MPI one-sided communication example

```
// rank 0 (origin)
MPI_Win window;
int buffer[SIZE] = ... ;
MPI_Win_create(&buffer, sizeof(int)*SIZE,
    sizeof(int), MPI_INFO_NULL,
    MPI_COMM_WORLD, &window);
MPI_Win_fence(MPI_MODE_NOSTORE |
    MPI_MODE_NOPUT | MPI_MODE_NOPRECEDE,
    window);
MPI_Put(&buffer, SIZE, MPI_INT, 1, 0,
    SIZE, MPI_INT, window);
MPI_Win_fence(MPI_MODE_NOSTORE |
    MPI_MODE_NOPUT | MPI_MODE_NOSUCCEED,
    window);
MPI_Win_free(&window);
```

```
// rank 1 (target)
MPI_Win window;
int buffer[SIZE] = { 0 };
MPI_Win_create(&buffer, sizeof(int)*SIZE,
    sizeof(int), MPI_INFO_NULL,
    MPI_COMM_WORLD, &window);
MPI_Win_fence(MPI_MODE_NOSTORE |
    MPI_MODE_NOPRECEDE | MPI_MODE_NOSUCCEED, window);
// window open, buffer can be written to
MPI_Win_fence(MPI_MODE_NOPUT |
    MPI_MODE_NOPRECEDE | MPI_MODE_NOSUCCEED, window);
// use buffer here
MPI_Win_free(&window);
```

# RMA semantics

---


- ▶ order of Get and Put/Accumulate is not guaranteed
  - ▶ race condition
  - ▶ same for multiple Put operations (use Accumulate instead)
- ▶ no local access to window during access epoch
  - ▶ use an RMA operation if absolutely required
- ▶ local vs. remote completion of operation
  - ▶ no send buffer re-use after Put until end of access epoch
- ▶ no concurrent passive synchronization epochs to same target
  - ▶ only relevant in multi-threading context
- ▶ lots of MPI fence optimizations
  - ▶ where to place, which assertions to use (start with 0 and add assertions as appropriate)

# Shared-memory one-sided communication


---

- ▶ one-sided communication also available in a truly shared memory fashion
  - ▶ origin of data transfer will get a pointer to access target memory
  - ▶ naturally only works for ranks in the same address space
  - ▶ c.f. POSIX shared memory segments
- ▶ allows to share memory between ranks
  - ▶ reduces memory footprint (e.g. no extra buffer for ghost cell exchange of a stencil)
  - ▶ even more efficient for intra-node communication than one-sided communication

# Ghost cell exchange (message passing in shared memory)



# Ghost cell exchange (MPI shared memory access)



## Example for shared memory one-sided communication

```
MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED,  
    0, MPI_INFO_NULL, &comm_sm);  
  
MPI_Win_allocate_shared((MPI_Aint)(sizeof(int) * 10),  
    sizeof(int), MPI_INFO_NULL, comm_sm, &recv_buf_ptr,  
    &win);  
  
MPI_Win_fence(0, win);  
*recv_buf_ptr = ...; // replaces the MPI_Put() call!  
MPI_Win_fence(0, win);  
MPI_Win_free(&win);
```

# MPI and multithreading

---

- ▶ one-sided shared memory communication can become quite complex
  - ▶ there are alternatives: MPI+(OpenMP or std::thread or Pthreads or TBB or ...)
  - ▶ these are known as *hybrid programming models*
  - ▶ both paradigms have their ups and downs (number of programming models, compiler support, compiler optimizations, interactions, side-effects, ...)
- ▶ MPI+threads needs to be supported by MPI, indicated by one of four safety levels
  - ▶ MPI\_THREAD\_SINGLE: only a single thread per rank
  - ▶ MPI\_THREAD\_FUNNELED: multithreaded ranks, but only the main thread calls MPI
  - ▶ MPI\_THREAD\_SERIALIZED: multithreaded, but only one per time calls MPI
  - ▶ MPI\_THREAD\_MULTIPLE: multithreaded, any thread can call MPI any time (with restrictions)
  - ▶ MPI implementations are not required to support more than MPI\_THREAD\_SINGLE
    - ▶ always check beforehand with MPI\_Query\_thread(...) and call MPI\_Init\_thread() instead of MPI\_Init()

## Error Handling

# Error handling

---

- ▶ MPI introduces additional hassle
  - ▶ imagine everything that can go wrong with a sequential process
  - ▶ add the fact that multiple processes are interacting, across the network
- ▶ default behavior: communication errors cause abort of MPI operation
  - ▶ causes respective process to exit
  - ▶ causes all other MPI processes of the same application to exit
- ▶ this only relates to halting crashes
  - ▶ also consider deadlocks or hangs

## Error handling cont'd

---

- ▶ all MPI routines return error codes
  - ▶ MPI\_SUCCESS if everything went well
- ▶ should always check and act accordingly
  - ▶ consider action to take when MPI calls fail
  - ▶ compare to a failing malloc in sequential program
  - ▶ make sure to free allocated resources (e.g. file handles)
  - ▶ inform the user!!

## Error handling cont'd

---

- ▶ error behavior can be altered with `MPI_Comm_set_errhandler(...)`
  - ▶ `MPI_ERRORS_ARE_FATAL`: abort if error detected (default)
  - ▶ `MPI_ERRORS_RETURN`: return error code to user program
- ▶ user-defined error handlers can be installed
  - ▶ `MPI_Comm_create_errhandler(...)`
  - ▶ `MPI_Comm_set_errhandler(...)`
  - ▶ `MPI_Comm_get_errhandler(...)`

# Additional topics

---

- ▶ **File I/O**
  - ▶ data file partitioning among processes, strided file access (derived data types!)
  - ▶ asynchronous data transfers
- ▶ **process management**
  - ▶ process spawning
  - ▶ socket-style communication
- ▶ **Fortran-specific issues**
  - ▶ slight differences in MPI function signatures
- ▶ **profiling support**
  - ▶ **MPI\_...** vs. **PMPI\_...**

# Summary

---

- ▶ **communicators and groups**
  - ▶ offers high-level control over sets of ranks
- ▶ **one-sided communication**
  - ▶ can improve performance, but can be tricky to use
- ▶ **error handling**
  - ▶ graceful exits